

Lecture 1

Paper to read: Map reduce

<https://static.googleusercontent.com/media/research.google.com/en//archive/mapreduce-osdi04.pdf>

Infrastructure for Applications

"We'd love to have abstractions"

Discussing about necessary infra needed for applications:

- Storage: (Accessing data abstracted away from application)
- Communication: (MapReduce is a specific example)
- Compute: (Networking, Reliability etc.)

The overall goal is to make distributed systems in such a way that they provide clean abstractions, so applications don't have to worry about underlying nature of implementation (k8s is good example).

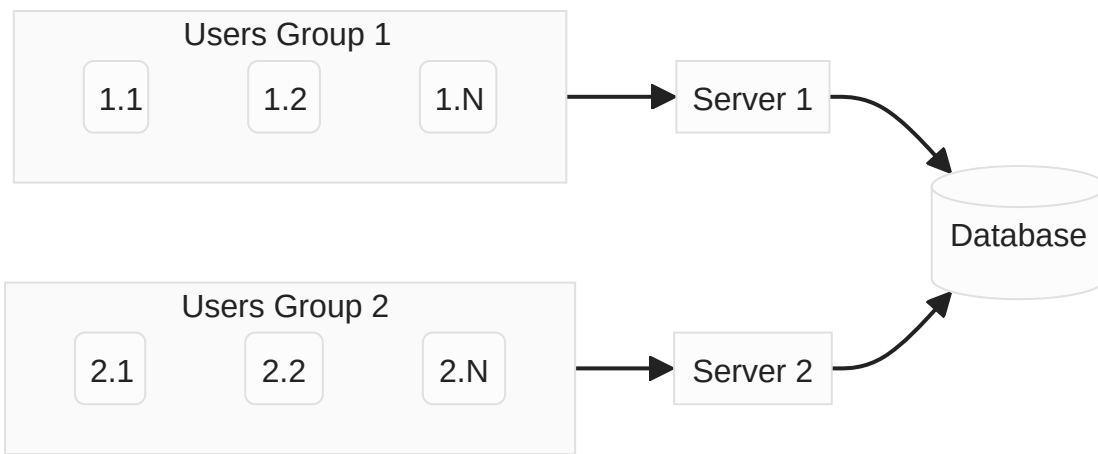
Common specific implementations include:

- RPCs
- Threads
- Concurrency (things about mutex, etc.)

Performance

"We are in general looking for scalable performance"

- Buy computers not programmers :)



This kind of scalability is basically theoretically.

So in this case adding more servers will not work since the bottleneck will become the Database.

Fault tolerance

"A system with a 1000s of computers there will always be a possibility that something will go wrong, that means one of the computers completely turning off"

Scale brings a lot of issues, since failures are common.

The overall idea is to hide/mask the failures that can occur in systems.

Main concepts under fault tolerance:

- Availability: make sure that application is at least available, replicasets are answer to some extent given not all of them are failing
- Recover-ability: in general, making sure that systems recover in case of any anomaly
 - Replication: the biggest issue here is that state of replicas might change leading to inconsistency
- Consistency:
 - Consider a simple service:
 - Get(k) -> v
 - Put(k, v)
 - Now when there are more than one copy of this service, you can have multiple copies of state.
 - Consider this
 - Initially all the RSs had Get(1) -> 20, we need to update to 21
 - Put(1, 21) (rs1)
 - Now we need to update rs2, but it fails to update (rs2)
 - Get(1), now rs2 will return stale copy

- To avoid such conditions we need to put some conditions (eg: eventual consistent systems)
- Strong/Weak consistency strategies
- Distributing service across regions to avoid natural disasters
 - But this leads to increased latency usually

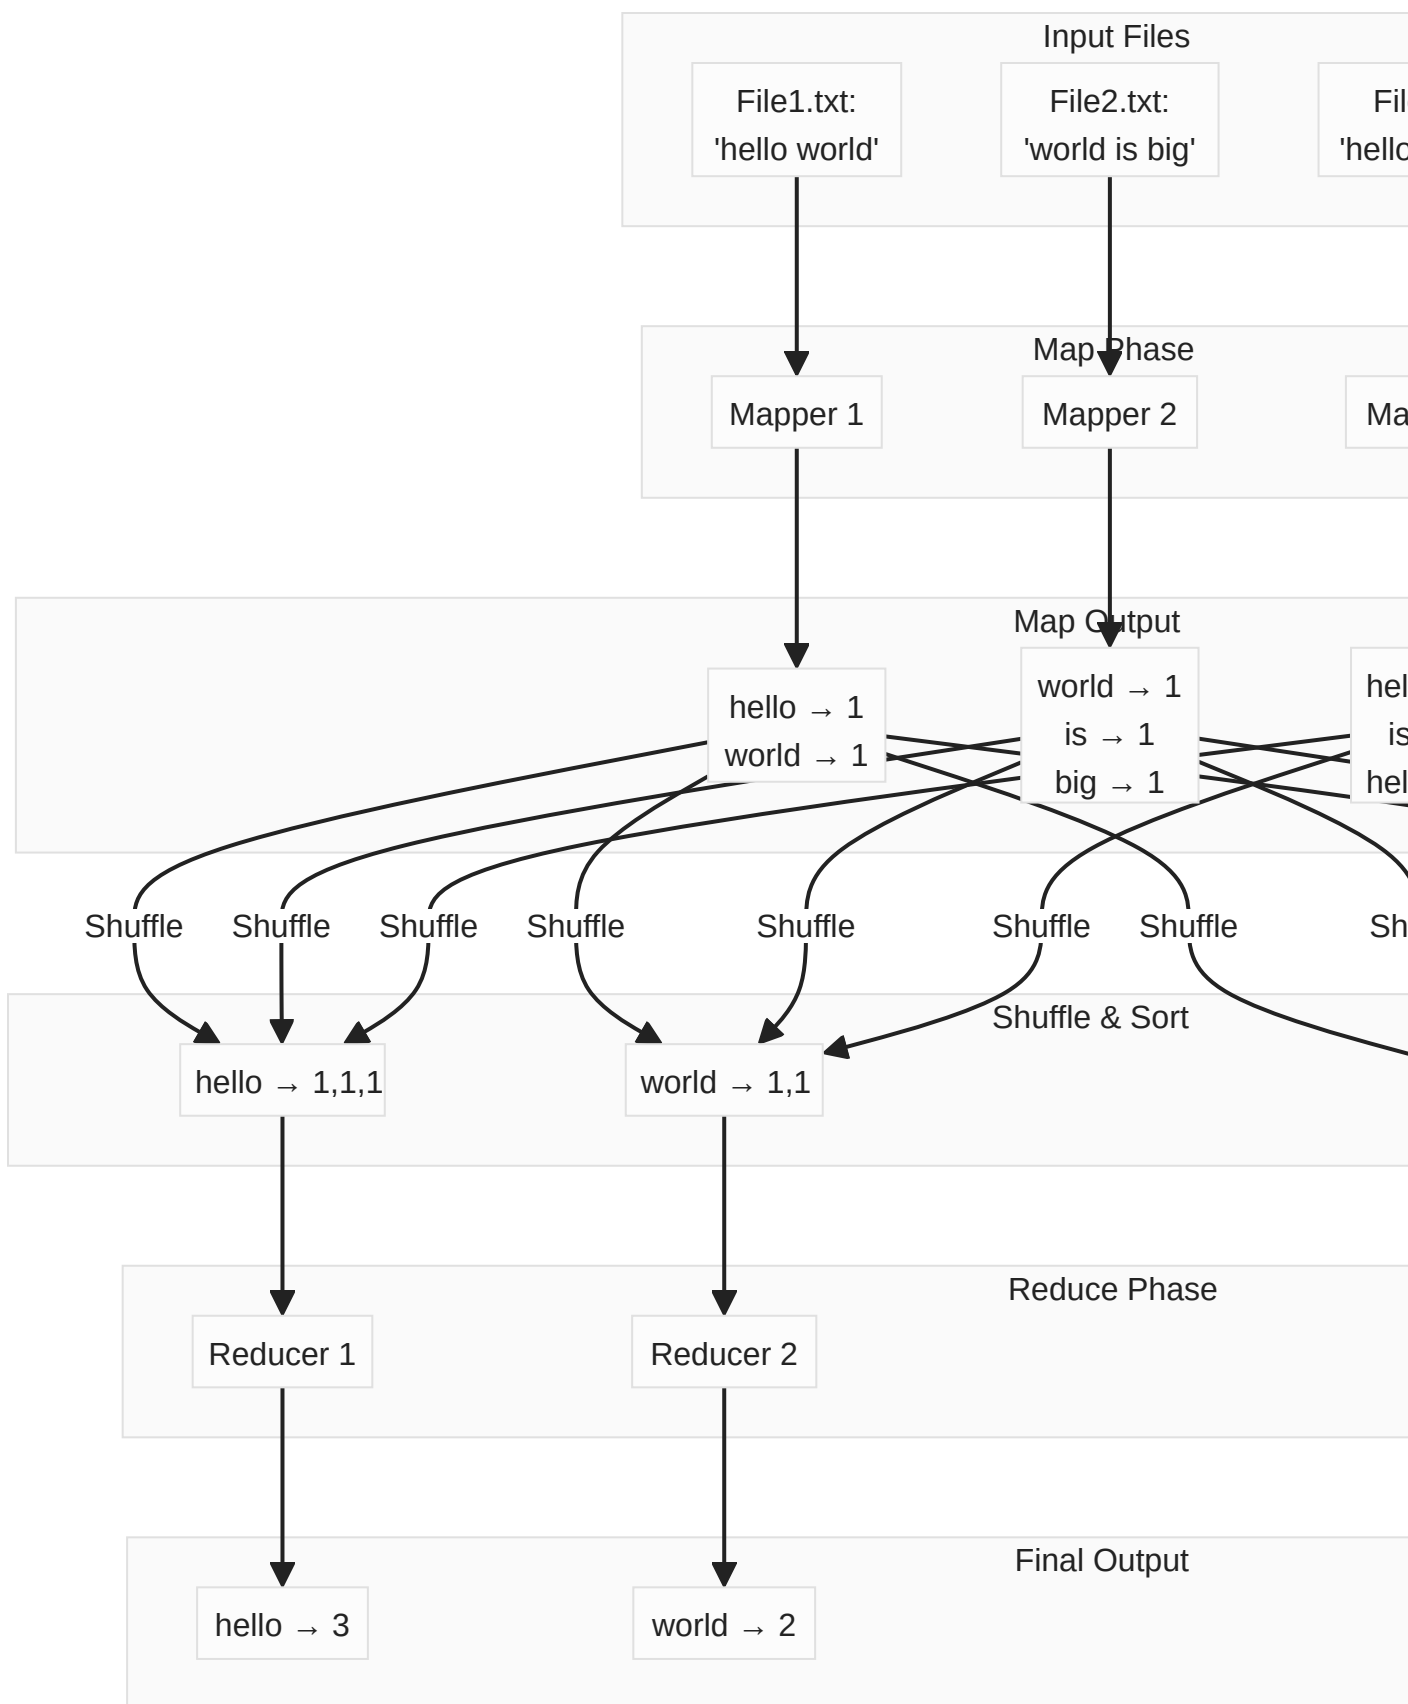
Case Study: MapReduce

Designed by Google

- Running huge computations at PetaByte scale (like indexing the whole fucking web lmao)
- Let engineers build and run applications without worrying about infrastructure (abstracting away using MapReduce)

```
Input: ArbitraryFileType
```

```
// Map function takes InputA and produces a key-val pair  
// Each key-val pair is then processed by a Reduce function  
// Map-Reduce together form a Task
```



An actual system that does map-reduce using Hadoop (distributing workers)

